

Reinforcement Learning

Building a Complete RL System

Filippos Christianos, Lukas Schäfer

11 February 2022



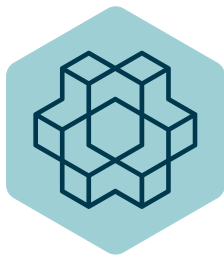
THE UNIVERSITY of EDINBURGH
informatics

Lecture Outline

- What is Gym?
- How to implement your own environment?
- How to implement a RL algorithm?
- How to evaluate your results?
- Demonstration

OpenAI Gym

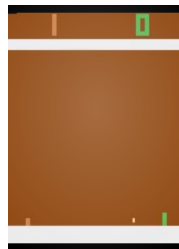

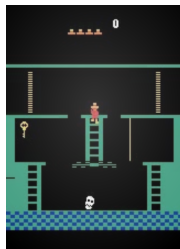
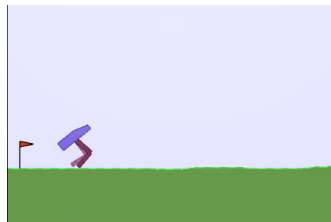
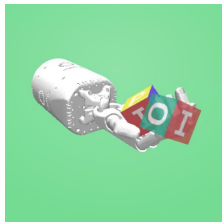
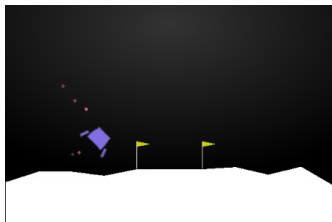
What is Gym? (Brockman et al., 2016)



- Open source interface for sequential decision processes
- Developed and maintained by OpenAI Research Lab
- Collection of RL environments
- Standardised interface for RL environments

```
pip install gym
```

Lots of Interesting Environments! (Brockman et al., 2016)



And many more... (Vinyals et al., 2017; Johnson et al., 2016; Kauten, 2018)



Gym Interface

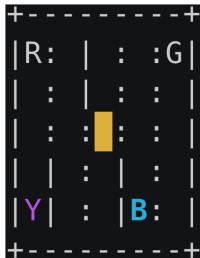
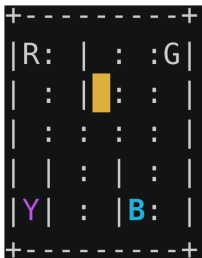
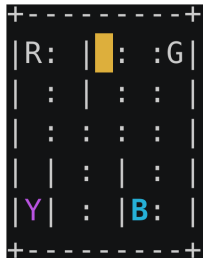
- `gym.make(<environment_name>)` → gym environment
Create a gym environment
- `env.reset()` → observation
Resets environment for a new episode
- `env.step(action)` → observation, reward, done, info
Take an action and observe new information
- `env.render()`
Render a visualisation of the current environmental state
- `env.close()`
Close created environment

Gym Example Snippet

Gym control flow



```
env = gym.make('CartPole-v0')
obs = env.reset()
done = False
while not done:
    env.render()
    action = agent.choose_action(obs)
    next_obs, reward, done, info = env.step(action)
    obs = next_obs
env.close()
```


Example: Taxi-v3 Environment



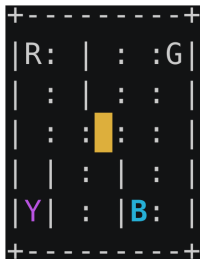
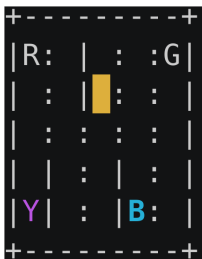
- Gridworld with 5×5 map
- R, G, Y, B - locations
 - B - passenger
 - Y - destination
 - taxi

Example: Taxi-v3 Environment



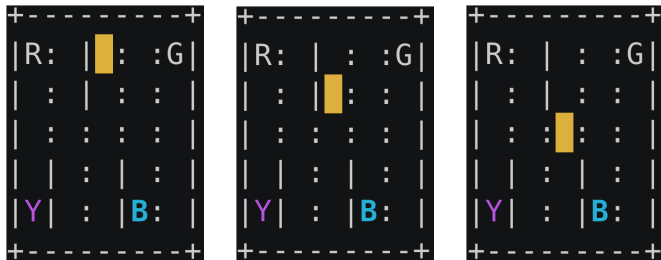
- Gridworld with 5×5 map
- R, G, Y, B - locations
B - passenger
Y - destination
■ - taxi
- Observations $\in [0, 499]$
including taxi row and col,
pass. and dest. index

Example: Taxi-v3 Environment



- Gridworld with 5×5 map
- R, G, Y, B - locations
B - passenger
Y - destination
■ - taxi
- Observations $\in [0, 499]$
including taxi row and col,
pass. and dest. index
- Actions:
South, North, East, West,
Pick, Drop

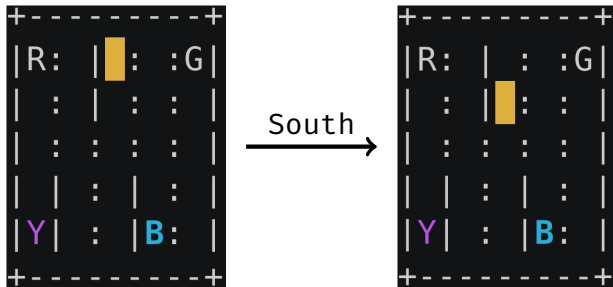
Example: Taxi-v3 Environment




- Goal: Pickup passenger and drop it off at destination
- Reward: +20 for successful delivery, -1 at each timestep, -10 for illegal move
- Challenge: navigate gridworld

- Gridworld with 5×5 map
- R, G, Y, B - locations
B - passenger
Y - destination
■ - taxi
- Observations $\in [0, 499]$
including taxi row and col,
pass. and dest. index
- Actions:
South, North, East, West,
Pick, Drop

Taxi Environment Step 1




Taxi Environment Step 1



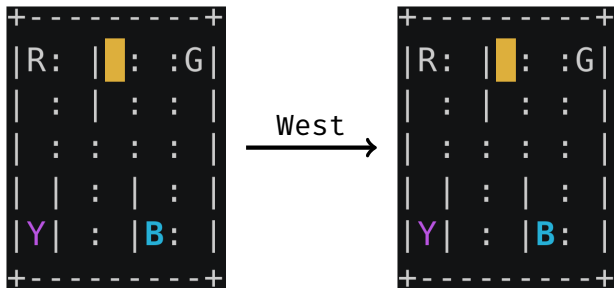
`nobs, r, done, _ = env.step(a):`

`o = 45` $\xrightarrow{a=0 \text{ (South)}}$ `<nobs = 154, r = -1, done = False>`

Taxi Environment Step II



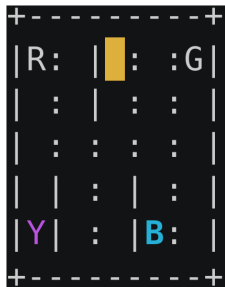
Taxi Environment Step II




```
nobs, r, done, _ = env.step(a):
```

```
o = 45  $\xrightarrow{a=3 \text{ (West)}}$  \langle nobs = 45, r = -1, done = False \rangle
```


Taxi Environment Step III



Taxi Environment Step III



```
nobs, r, done, _ = env.step(a):
```

```
o = 45  $\xrightarrow{a=4/5 \text{ (Pick/ Drop)}}$  \langle nobs = 45, r = -10, done = False \rangle
```

Implement your RL Agent

Recap: SARSA

On-Policy TD Control: Sarsa

→ learn q_π and improve π while following π

Updates: $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$

Exploration: ϵ -greedy policy π

Recap: SARSA

On-Policy TD Control: Sarsa

→ learn q_π and improve π while following π

Updates: $Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$

Exploration: ϵ -greedy policy π

Initialize $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$, arbitrarily, and $Q(\text{terminal-state}, \cdot) = 0$

Repeat (for each episode):

 Initialize S

 Choose A from S using policy derived from Q (e.g., ϵ -greedy)

 Repeat (for each step of episode):

 Take action A , observe R, S'

 Choose A' from S' using policy derived from Q (e.g., ϵ -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

 until S is terminal

SARSA Agent Class Structure

- **__init__**: Initialise agent and Q-table as dictionary mapping (obs, act) -> q-val
- **act**: ϵ -greedy policy
- **learn**: Update Q-table given new experience

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

- **schedule_hyperparameters**: Update hyperparameters given training progress

And now in Code ... act

Epsilon-greedy Action Selection

```
def act(self, obs):  
    act_vals = [self.q_table[(obs, act)] for act in range(self.  
n_acts)]  
    max_val = max(act_vals)  
    max_acts = [idx for idx, act_val in enumerate(act_vals) if  
act_val == max_val]  
  
    if random.random() < self.epsilon:  
        return random.randint(0, self.n_acts - 1)  
    else:  
        return random.choice(max_acts)
```

And now in Code ... learn

SARSA Q-Update


```
def learn(self, obs, action, reward, n_obs, n_action, done):  
    target_value = reward + self.gamma * (1 - done) * self.  
    q_table[(n_obs, n_action)]  
    self.q_table[(obs, action)] += self.alpha * (  
        target_value - self.q_table[(obs, action)]  
    )  
    return self.q_table[(obs, action)]
```

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha [R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t)]$$

And now in Code ... schedule_hyperparameters

SARSA ϵ -Scheduling

```
def schedule_hyperparameters(self, timestep, max_timestep):  
    max_deduct, decay = 0.95, 0.07  
    self.epsilon = 1.0 - (min(1.0, timestep/(decay *  
    max_timestep))) * max_deduct
```



Evaluate your Results

Why do We Evaluate in the First Place?

- It gives our approach credibility
- Empirical evaluation is no proof, but can give strong indication about the strengths and limitations of an approach (when done right!)

Why do We Evaluate in the First Place?

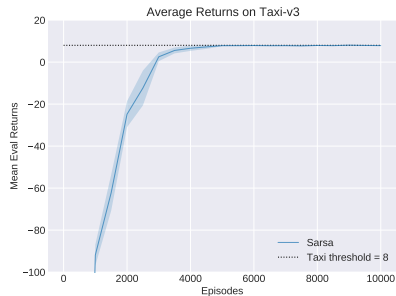
- It gives our approach credibility
- Empirical evaluation is no proof, but can give strong indication about the strengths and limitations of an approach (when done right!)

How to do it *right*?

What to Evaluate?

Evaluation Returns

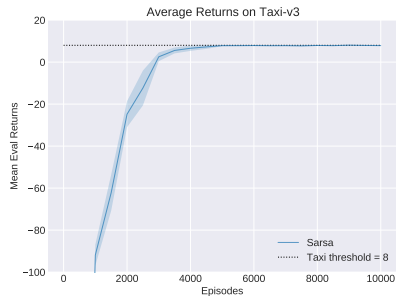
- Plot mean returns over multiple runs
- Visualise standard deviation or confidence interval



What to Evaluate?

Evaluation Returns

- Plot mean returns over multiple runs
- Visualise standard deviation or confidence interval

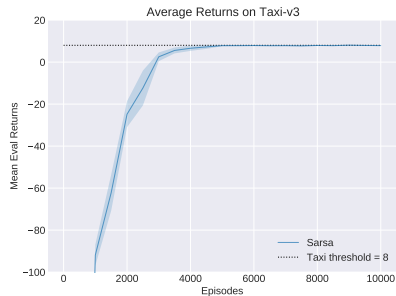


Which returns do we plot?

What to Evaluate?

Evaluation Returns

- Plot mean returns over multiple runs
- Visualise standard deviation or confidence interval





Which returns do we plot?

- Execute multiple evaluation runs with $\epsilon = 0$ at fixed intervals
- Evaluation does not involve any learning!

Keep Track of Everything!

Hyperparameters

- Track hyperparameters, here ϵ -decay
- Try various values in a grid- or random-search and find good configuration



SARSA Gridsearch over Learning Rate α for Taxi-v3




Figure 1: $\alpha = 0.9$

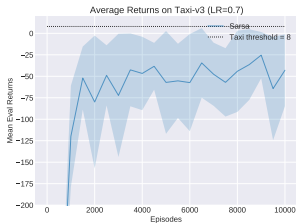


Figure 3: $\alpha = 0.7$

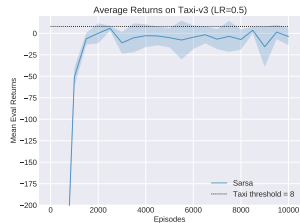


Figure 5: $\alpha = 0.5$




Figure 2: $\alpha = 0.3$

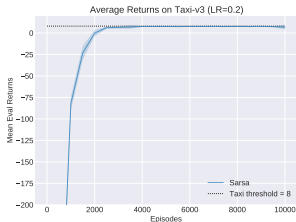


Figure 4: $\alpha = 0.2$

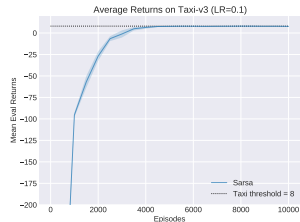


Figure 6: $\alpha = 0.1$

SARSA Learning Rate α Gridsearch Overview




Figure 7: Gridsearch overview over learning rate α with half standard deviation as shading

"But it worked last time!"

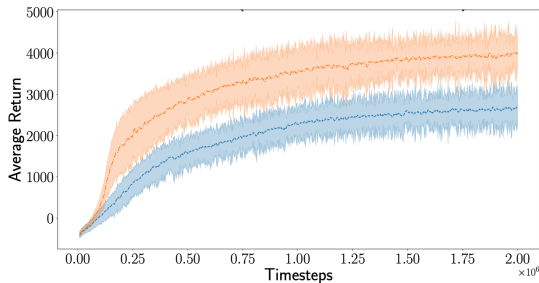
”But it worked last time!”

- It's not enough to make it work once!
- Meaningful evaluation achieves consistent performance over multiple randomised runs
- Most RL algorithms have random components (e.g. ϵ -greedy policies)

Is plotting the mean return, even with confidence interval, enough?

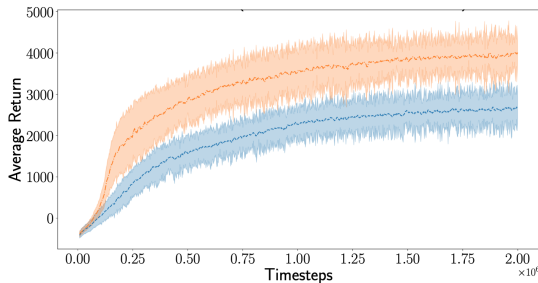
Common Pifalls (2) (Henderson, 2018; Colas et al., 2019)

Is plotting the mean return, even with confidence interval, enough?



Common Pifalls (2) (Henderson, 2018; Colas et al., 2019)


Is plotting the mean return, even with confidence interval, enough?



Which one is better?

Common Pifalls (2) (Henderson, 2018; Colas et al., 2019)


Is plotting the mean return, even with confidence interval, enough?



Which one is better? It's actually the same method!

Common Pifalls (2) (Henderson, 2018; Colas et al., 2019)

Is plotting the mean return, even with confidence interval, enough?



Which one is better? It's actually the same method!

Apparently, it's not enough! \rightarrow Statistical hypothesis testing (Colas et al., 2019)
and effective statistical evaluation (Agarwal et al., 2021)

"Why should I use those random seeds? `random` already delivers random values!"

Common Pitfalls (3) (Bishop, 2019)

”Why should I use those random seeds? **random** already delivers random values!”

- Our goal with empirical evaluations is to make meaningful claims about the implemented approach and achieve **reproducible** performance
- Random seeds allow us to fixate random behaviour
- Reproducibility is key for meaningful research

Common Pitfalls (3) (Bishop, 2019)

”Why should I use those random seeds? **random** already delivers random values!”

- Our goal with empirical evaluations is to make meaningful claims about the implemented approach and achieve **reproducible** performance
- Random seeds allow us to fixate random behaviour
- Reproducibility is key for meaningful research

But NEVER choose/ tune your random seeds!

“Why should I use those random seeds? **random** already delivers random values!”

- Our goal with empirical evaluations is to make meaningful claims about the implemented approach and achieve **reproducible** performance
- Random seeds allow us to fixate random behaviour
- Reproducibility is key for meaningful research

But NEVER choose/ tune your random seeds!

Rein in the four horsemen of irreproducibility



Dorothy Bishop describes how threats to reproducibility, recognized but unaddressed for decades, might finally be brought under control.

Demonstration

All code is available at [https://github.com/uoε-agents/
Building-a-Complete-RL-System_Demonstration](https://github.com/uoε-agents/Building-a-Complete-RL-System_Demonstration)

References

- Agarwal, R., Schwarzer, M., Castro, P. S., Courville, A. C., and Bellemare, M. (2021). Deep reinforcement learning at the edge of the statistical precipice. In *Advances in Neural Information Processing Systems*, volume 34.
- Bishop, D. (2019). Rein in the Four Horsemen of Irreproducibility. *Nature*, 568(7753).
- Brockman, G., Cheung, V., Pettersson, L., Schneider, J., Schulman, J., Tang, J., and Zaremba, W. (2016). OpenAI Gym.

- Colas, C., Sigaud, O., and Oudeyer, P.-Y. (2019). A Hitchhiker's Guide to Statistical Comparisons of Reinforcement Learning Algorithms. *arXiv preprint arXiv:1904.06979*.
- Henderson, P. (2018). *Reproducibility and Reusability in Deep Reinforcement Learning*. PhD thesis, McGill University Libraries.
- Johnson, M., Hofmann, K., Hutton, T., and Bignell, D. (2016). The Malmo Platform for Artificial Intelligence Experimentation. In *IJCAI*, pages 4246–4247.
- Kauten, C. (2018). Super Mario Bros for OpenAI Gym. GitHub.
- Vinyals, O., Ewalds, T., Bartunov, S., Georgiev, P., Vezhnevets, A. S., Yeo, M., Makhzani, A., Küttler, H., Agapiou, J., Schrittwieser, J., et al. (2017). Starcraft II: A new Challenge for Reinforcement Learning. *arXiv preprint arXiv:1708.04782*.

Any questions about this lecture or the demonstration?